CS-497, Spring 2024
Homework #3: Question Answering with Transformers
Due Date: Monday, May 20th @ 11:59PM
Total Points: 10.0

Description: Question-answering with transformer-based architectures is the dominant paradigm in natural language processing. Among the various formulations of this task, multiple-choice question-answering is among the most tractable. In this assignment, you will perform question-answering on the OpenBookQA dataset (see https://arxiv.org/abs/1809.02789). OpenBookQA is a multiple-choice question-answering dataset of elementary school-level scientific questions. Questions are posed in natural language with four possible choices. Background knowledge in the form of a one-sentence "fact" is also provided.

This task's primary formulation requires you to identify the correct "fact" from among 1,326 facts. To focus on transformer architectures in this assignment, you are given the "fact" associated with each question. The training dataset consists of approximately 5,000 observations; the validation and test sets have 500 observations each. Each observation contains a question stem, the associated fact, four labeled choices and a correct answer label. You may ignore all other information.

You will use transformer-based architecture to perform multiple-choice question-answering for classification and generative approaches. Starter code is provided to read the JSON dataset files. You may need to install Huggingface Transformers (see https://huggingface.co/docs/transformers/installation) on your machine.

Tasks:

1. **Classification Approach (5 pts):** For this approach, you must use the bert-base-uncased checkpoint for the BertModel (see https://huggingface.co/docs/transformers/v4.26.1/en/model_doc/bert) and fine-tune the model on the training set using your own code. Implementations using BertForMultipleChoice or other specialized variants will not be accepted. The format of the text input to BERT and classification methodology for this task is up to you. One approach is to leverages the `[CLS]` token, that can be implemented as:

   a. Encode each instance as
   ```
   [CLS] <fact> <stem> <text of choice #1> [END]
   [CLS] <fact> <stem> <text of choice #2> [END]
   [CLS] <fact> <stem> <text of choice #3> [END]
   [CLS] <fact> <stem> <text of choice #4> [END]
   ```
   b. Retrieve the context embedding at the top of the transformer stack for the `[CLS]` token and take the dot product with a `[d_mod x 1]` linear layer.
   c. Treat the resulting vector as logits, apply the softmax and compute a multiclass cross-entropy loss over the four choices.
   d. Fine-tune BERT the complete dataset with this encoding for multiple epochs.
   e. Use the trained model to perform inference on the validation and test sets.
   f. Report accuracy

   Please describe your approach sufficiently so that someone familiar with transformers can replicate your results. Also, report your zero-shot and fine-tuned accuracies on the validation and test sets. Please discuss any limitations of your approach and possible solutions. Using this approach described, I was able to achieve an accuracy of 55.9% on the test set.

2. **Generative Approach (5 pts):** For this approach, you are free to select a decoder-only or an encoder-decoder architecture of your choice. You must design a prompt and output format to perform this task. The design will include a methodology of interpreting generated output to evaluate model performance (which may consist of exact match, ROUGE-F1 scores, etc.) and must accurately capture model performance. You may fine-tune your architecture using the scripts in transformers/examples and use functions like `model.generate()` (see https://huggingface.co/blog/how-to-generate). One approach is to:

   a. Encode each instance as `[START] <fact> <stem> [A] <text of choice #1> [B] <text of choice #2> [C] <text of choice #3> [D] <text of choice #4> [ANSWER] <correct label A, B, C, or D>`
   b. Fine-tune your model on the complete dataset with this encoding for multiple epochs. (Note: Using the loss on the last token only may boost performance.)
   c. Use the trained model to generate a label following the [ANSWER] token on the validation and test.
   d. Use exact match to calculate the accuracy of your approach.

   Please describe your approach in sufficient detail such that someone familiar with transformers can replicate your results. Also, report your zero-shot and fine-tuned accuracies on the validation and test sets, and provide examples of your prompts and outputs. Please discuss any limitations of your approach and possible solutions. Discuss the performance of your generative approach compared to your classification approach. Using an exact match performance metric based on A, B, C or D encoding, I can achieve an accuracy of 40.8% on the test set using GPT-2 (https://huggingface.co/gpt2#training-procedure) and my own code for fine tuning. Another approach is

   a. Encode each instance as `[START] <fact> <stem> [ANSWER] <text of correct answer>`
   b. Fine-tune your model on the complete dataset with this encoding for multiple epochs.
   c. Use the trained model to generate answers on the validation and test sets using a decoding method of your choice.
   d. Use an automatic metric like ROUGE-F1 treating the answer choices as the reference text.
   e. Select the multiple choice answer as the one with the highest ROUGE-F1 score.
   f. Report accuracy.

   Using this I was able to achieve an accuracy of only 27.6%.

What to Submit: Please submit a single .zip file for the group. The zip file should contain the following:

- A single .pdf for all written responses and results.
- Individual .py files (two files maximum) for your Classification and Generative approaches.

Only one group member needs to submit.